

Theo Farrell

theo.farrell99@outlook.com | <https://theosdoor.github.io>

Last updated: 05/10/2026

AI safety researcher and field-builder. First-class MSci Natural Sciences (Computer Science & Philosophy) from Durham University. First-author on two interpretability papers (NeurIPS 2025, ICML 2026) and contributor to LLM-safety work at IASEAI 2026. Currently a Research Fellow at LASR Labs working on the science of model organisms of misalignment. Founded Durham's AI Safety Initiative and secured \$19,000 in research and organising grants.

EDUCATION

Durham University

2022 – 2026

MSci Natural Sciences (Computer Science & Philosophy)

First Class Honours (Class I, 79%)

- **Dissertation (82%):** “Relational Composition and Sparse Autoencoders”, supervised by Noura Al Moubayed.
- **Key Modules:** NLP, Deep Learning, Computer Vision, Reinforcement Learning.
- **Annual End-of-Year Averages:** Y1: 76%, Y2: 76.5%, Y3: 77%, Y4: 82%.
- **Top Module Rankings:** Natural Computing (1st/51), Computer Vision (2nd/81), Moral Theory (1st/135), Political Philosophy (2nd/138).

EXPERIENCE

LASR Labs

07/2026 – Present

Research Fellow, London (mentored by Satvik Golechha)

- Full-time AI safety research programme in London.
- Currently working on science of model organisms by devising evaluations for measuring how well they exhibit an implanted behaviour.

Founder and Lead Organiser

09/2023 – 02/2026

Durham AI Safety Initiative

- Organised recruitment and increased our weekly attendance from 5 to 20 participants.
- Facilitated BlueDot discussion groups to upskill members in AI alignment and policy knowledge.
- Led technical upskilling workshops in deep learning and NLP based on ARENA.
- Pathfinder Fellow since Dec 2024 - received training and mentorship for AI Safety organising after a competitive application process.

University of Calgary

06/2024 – 09/2024

MITACS Global Research Intern (supervised by Aditya Nittala)

- Built soft electrotactile haptic devices for human-computer interaction. Contributed to a paper at ACM UIST'24.

RELEVANT PUBLICATIONS [Scholar]

- Farrell, Theo, Patrick Leask, and Noura Al Moubayed. “Order by Scale: Relative-Magnitude Relational Composition in Attention-Only Transformers.” Socially Responsible and Trustworthy Foundation Models at NeurIPS 2025. Link: <https://openreview.net/forum?id=vWRVzNtk7W>
- Farrell, Theo, Patrick Leask, and Noura Al Moubayed. “Sparse Autoencoders Can Learn Graded Latents for Relational Composition.” Mechanistic Interpretability Workshop at ICML 2026. Link: <https://openreview.net/forum?id=ltQ6XAduTF>
- Kempermann, Manon, et al. “Safe for whom? Rethinking how we evaluate the safety of LLMs for real users.” Second Conference of the International Association for Safe and Ethical Artificial Intelligence (IASEAI'26), 2026. Link: <https://ojs.aaai.org/index.php/IASEAI/article/view/43033>

Reviewing Activities: ICML 2026 Mechanistic Interpretability Workshop, NeurIPS 2025 Mechanistic Interpretability Workshop, NeurIPS 2025 ResponsibleFM Workshop.

GRANTS AND AWARDS

- **\$4,500 from Pathfinder** (10/2025): 1 year of group expenses for Durham AI Safety Initiative.
- **\$6,500 from Coefficient Giving** (07/2025): 4 months of independent mechanistic interpretability research.
- **\$8,000 from Coefficient Giving** (03/2025): 1 year of group expenses for Durham AI Safety Initiative.
- **L4 Natural Sciences Outstanding Achievement** (07/2026): Averaged over 75% in 4th year.
- **L3 Natural Sciences Outstanding Achievement** (07/2025): Averaged over 75% in 3rd year.
- **L2 Natural Sciences Outstanding Achievement** (07/2024): Averaged over 75% in 2nd year.

CERTIFICATIONS

- **AGI Strategy** (BlueDot Impact, 2026). [[Certificate](#)]

RELEVANT TECHNICAL SKILLS

Programming Python, C#, Javascript, HTML.

Tools PyTorch, TransformerLens, SAELens, Hugging Face, NumPy, scikit-learn, Weights & Biases, pandas, Git, uv, $\text{\LaTeX}$.